

Loan default prediction

January 12, 2021

1 Loan default Detection

We will be investigating a data set from a bank that has customer information such as ‘Customer Id’, ‘Age’, ‘Edu’, ‘Years Employed’, ‘Income’, ‘Card Debt’, ‘Other Debt’, ‘Address’, ‘DebtIncomeRatio’. An additional feature in the data set is “Defaulted” which is our target data. If Defaulted is equal to 0 it means that the customer has paid back the loan and if it is equal to 1 it means that the customer has defaulted on the loan.

Many customers are still paying back their loans and we will be predicting whether they will default or not based on the customer information using a suitable classification model.

The used data set is publicly available and can be downloaded via: https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/ML0101ENv3/labs/Cust_Segmentation.csv

1.1 1- Loading the Data

```
[28]: import pandas as pd
import matplotlib
from matplotlib import pyplot as plt
import seaborn as sns
import numpy as np
```

```
[29]: df = pd.read_csv("C:/Users/eliec/Downloads/customer_segmentation.csv")
#Data source shorturl.at/dhFW7
```

```
[30]: pd.options.display.max_rows = 20
pd.options.display.max_columns = 100
print(df.shape)
df.head()
```

(850, 10)

```
[30]:
```

	Customer Id	Age	Edu	Years Employed	Income	Card Debt	Other Debt	\
0	1	41	2	6	19	0.124	1.073	
1	2	47	1	26	100	4.582	8.218	
2	3	33	2	10	57	6.111	5.802	
3	4	29	2	4	19	0.681	0.516	
4	5	47	1	31	253	9.308	8.908	

	Defaulted	Address	DebtIncomeRatio
0	0.0	NBA001	6.3
1	0.0	NBA021	12.8
2	1.0	NBA013	20.9
3	0.0	NBA009	6.3
4	0.0	NBA008	7.2

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 850 entries, 0 to 849
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Customer Id           850 non-null   int64
1   Age                   850 non-null   int64
2   Edu                   850 non-null   int64
3   Years Employed        850 non-null   int64
4   Income                850 non-null   int64
5   Card Debt             850 non-null   float64
6   Other Debt            850 non-null   float64
7   Defaulted             700 non-null   float64
8   Address               850 non-null   object
9   DebtIncomeRatio       850 non-null   float64
dtypes: float64(4), int64(5), object(1)
memory usage: 66.5+ KB
```

our target variable “Defaulted” has 150 entries missing out of 850 which are the entries that we need to predict. We will first split the complete data that has all the entries (700) into test and train data to be able to identify the best classification model to use. Once we have identified the best model we will proceed to predicting the outcome of the loans that are still outstanding.

1.2 2- Data Wrangling

```
[32]: # Slicing the data where teh feature "Defaulted" is missing. This represents
      ↪ the oustanding loans
df_missing = df[df["Defaulted"].isnull()]
df_missing.head()
```

```
[32]:   Customer Id  Age  Edu  Years Employed  Income  Card Debt  Other Debt  \
5           6   40    1           23      81      0.998      7.831
8           9   26    1           5      18      0.575      2.215
11          12   34    2           9      40      0.374      0.266
13          14   46    1           6      30      1.415      3.865
15          16   24    1           1      16      0.185      1.287
```

Defaulted	Address	DebtIncomeRatio
-----------	---------	-----------------

5	NaN	NBA016	10.9
8	NaN	NBA006	15.5
11	NaN	NBA003	1.6
13	NaN	NBA019	17.6
15	NaN	NBA005	9.2

```
[33]: # dropping the rows with missing "Defaulted" from the dataframe
```

```
df = df.dropna()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 700 entries, 0 to 849
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Customer Id           700 non-null   int64
1   Age                   700 non-null   int64
2   Edu                   700 non-null   int64
3   Years Employed        700 non-null   int64
4   Income                700 non-null   int64
5   Card Debt             700 non-null   float64
6   Other Debt            700 non-null   float64
7   Defaulted             700 non-null   float64
8   Address               700 non-null   object
9   DebtIncomeRatio       700 non-null   float64
dtypes: float64(4), int64(5), object(1)
memory usage: 60.2+ KB
```

1.3 3- Feature Analysis

```
[56]: print(len(df["Address"].unique()))
print(len(df["Customer Id"].unique()))
```

```
31
700
```

The feature “Address” has 31 unique values even though we have 700 different customers which means that the addresses represent neighborhoods or cities. Sometimes neighborhoods have a strong correlation to socio-economic factors that might have an effect on our target values. We will investigate that correlation but first we need to cast the Addresses from object to numbers

```
[57]: df["Address"] = df["Address"].astype('category')
df["Address_code"] = df["Address"].cat.codes
```

```
[58]: # checking the correlation between the features and our target variable
```

```
df.corr()["Defaulted"].abs().sort_values(ascending=False)
```

```
[58]: Defaulted          1.000000
      DebtIncomeRatio    0.389575
      Years Employed     0.282978
      Card Debt          0.244734
      Address_code       0.164832
      Other Debt         0.145716
      Age                0.137657
      Edu                0.114676
      Income             0.070970
      Customer Id        0.004244
      Name: Defaulted, dtype: float64
```

```
[44]: # checking the correlation between all features
```

```
fig = plt.figure(figsize=(10,8))
corr = df.corr().abs()
sns.heatmap(corr)
df.corr().abs()
```


```
[44]:
```

	Customer Id	Age	Edu	Years Employed	Income \
Customer Id	1.000000	0.014894	0.039896	0.034290	0.050366
Age	0.014894	1.000000	0.022325	0.536497	0.478710
Edu	0.039896	0.022325	1.000000	0.153621	0.235190
Years Employed	0.034290	0.536497	0.153621	1.000000	0.619681
Income	0.050366	0.478710	0.235190	0.619681	1.000000
Card Debt	0.018939	0.295214	0.088277	0.403698	0.570196
Other Debt	0.031602	0.340213	0.165458	0.406089	0.610663
Defaulted	0.004244	0.137657	0.114676	0.282978	0.070970
DebtIncomeRatio	0.002401	0.016398	0.008838	0.031182	0.026777
Address_code	0.057768	0.597081	0.057163	0.322342	0.316326

	Card Debt	Other Debt	Defaulted	DebtIncomeRatio \
Customer Id	0.018939	0.031602	0.004244	0.002401
Age	0.295214	0.340213	0.137657	0.016398
Edu	0.088277	0.165458	0.114676	0.008838
Years Employed	0.403698	0.406089	0.282978	0.031182
Income	0.570196	0.610663	0.070970	0.026777
Card Debt	1.000000	0.633108	0.244734	0.501772
Other Debt	0.633108	1.000000	0.145716	0.584867
Defaulted	0.244734	0.145716	1.000000	0.389575
DebtIncomeRatio	0.501772	0.584867	0.389575	1.000000
Address_code	0.208779	0.227462	0.164832	0.012248

	Address_code
Customer Id	0.057768
Age	0.597081
Edu	0.057163

Years Employed	0.322342
Income	0.316326
Card Debt	0.208779
Other Debt	0.227462
Defaulted	0.164832
DebtIncomeRatio	0.012248
Address_code	1.000000



We can see that many features are correlated to one another especially Age and Years Employed so we will not include Age since it is already reflected by YearsEmployed. We will not include Income in the feature selection as it has a low correlation with our target and is represented by a more important feature which is DebtIncomeRatio which has the highest correlation with our target.

The address_code has a weak correlation with our target value but also has high correlation with other features such as Age, Income and Debt so we will not include it in the feature selection

1.4 4- Model Deployment and Evaluation

We will evaluate the following models and choose the one with the highest accuracy

- Decision Tree
- Random Forest
- Logistic Regression
- Support Vector Machine
- K Nearest Neighbors
- Gradient Boosting

Splitting and preprocessing the data

```
[59]: # splitting the data into test and train

from sklearn.model_selection import train_test_split
X = df[['Edu', 'Years Employed', 'Card Debt', 'Other Debt', 'DebtIncomeRatio']]
y = df["Defaulted"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↪random_state=1, stratify=y)
```

Logistic Regression, SVM and KNN require feature standardization so we will do that for all the models

```
[60]: from sklearn.preprocessing import StandardScaler
sc= StandardScaler()
sc.fit(X_train)
X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)
```

Deploying the models

```
[61]: from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
```

```
[62]: forest = RandomForestClassifier()
forest.fit(X_train_std, y_train)
LR = LogisticRegression()
LR.fit(X_train_std, y_train)
svm = SVC()
svm.fit(X_train_std, y_train)
tree = DecisionTreeClassifier()
tree.fit(X_train_std, y_train)
knn = KNeighborsClassifier()
knn.fit(X_train_std, y_train)
gbc = GradientBoostingClassifier()
```

```
gbc.fit(X_train_std, y_train)
```

```
[62]: GradientBoostingClassifier()
```

Evaluationg the models

```
[64]: y_hat_forest = forest.predict(X_test_std)
y_hat_LR = LR.predict(X_test_std)
y_hat_svm = svm.predict(X_test_std)
y_hat_tree = tree.predict(X_test_std)
y_hat_knn = knn.predict(X_test_std)
y_hat_gbc = gbc.predict(X_test_std)
```

```
[69]: def get_name(list_):
    name = [x for x in globals() if globals()[x] is list_][0]
    return name
```

```
[70]: models = [y_hat_forest, y_hat_LR, y_hat_svm, y_hat_tree, y_hat_knn, y_hat_gbc]
```

```
[74]: from sklearn.metrics import accuracy_score
for model in models:
    print (get_name(model), "accuracy: ", accuracy_score(y_test, model))
```

```
y_hat_forest accuracy: 0.7485714285714286
y_hat_LR accuracy: 0.8228571428571428
y_hat_svm accuracy: 0.8
y_hat_tree accuracy: 0.6914285714285714
y_hat_knn accuracy: 0.7314285714285714
y_hat_gbc accuracy: 0.7657142857142857
```

Without tuning any Hyperparameters, it seems like the LogisticRegression is the superior model for predicting loan default. It is also the most convenient because it also predicts the probability of defaulting as well.

2 5- Predicting the outcome of the outstanding loans using Logistic Regression

```
[26]: X_missing = df_missing[['Edu', 'Years Employed', 'Card Debt', 'Other Debt', '
    ↳ 'DebtIncomeRatio']]
sc.fit(X)
X_std = sc.transform(X)
X_missing_std = sc.transform(X_missing)
LR.fit(X_std, y)
y_hat_missing = LR.predict(X_missing_std)
```

```
[77]: # Adding the values of "y_hat_missing" to the column "Defaulted" for the
    ↳ outstanding loans
```

```
df_missing["Defaulted"] = y_hat_missing
df_missing
```

```
[77]:
```

	Customer	Id	Age	Edu	Years	Employed	Income	Card Debt	Other Debt	\
5		6	40	1		23	81	0.998	7.831	
8		9	26	1		5	18	0.575	2.215	
11		12	34	2		9	40	0.374	0.266	
13		14	46	1		6	30	1.415	3.865	
15		16	24	1		1	16	0.185	1.287	
..			
818		819	35	2		0	35	2.383	1.957	
820		821	37	1		4	24	0.419	2.989	
825		826	32	2		12	116	4.027	2.585	
835		836	21	3		0	41	2.367	5.628	
845		846	27	1		5	26	0.548	1.220	

	Defaulted	Address	DebtIncomeRatio
5	0.0	NBA016	10.9
8	0.0	NBA006	15.5
11	0.0	NBA003	1.6
13	0.0	NBA019	17.6
15	0.0	NBA005	9.2
..
818	1.0	NBA006	12.4
820	0.0	NBA010	14.2
825	0.0	NBA011	5.7
835	1.0	NBA001	19.5
845	0.0	NBA007	6.8

[150 rows x 10 columns]

```
[ ]:
```